

data-linear-regression

April 20, 2025

1 Linear Regression Model

1.1 Import Essential Libraries

```
[7]: # Basic Libraries
import numpy as np
import pandas as pd
import seaborn as sb
import matplotlib.pyplot as plt # we only need pyplot
sb.set() # set the default Seaborn style for graphics
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import mean_squared_error, r2_score, mean_absolute_error
```

1.2 Import Data

```
[5]: df = pd.read_csv('datasets/twitch-data-cleaned.csv')
df.head()
```

```
[5]:
```

	channel	watch_time_minutes	stream_time_minutes	peak_viewers	\
0	xQcOW	6196161750	215250	222720	
1	summit1g	6091677300	211845	310998	
2	Gaules	5644590915	515280	387315	
3	ESL_CSGO	3970318140	517740	300575	
4	Tfue	3671000070	123660	285644	

	average_viewers	followers	followers_gained	views_gained	partnered	\
0	27716	3246298	1734810	93036735	True	
1	25610	5310163	1370184	89705964	True	
2	10976	1767635	1023779	102611607	True	
3	7714	3944850	703986	106546942	True	
4	29602	8938903	2068424	78998587	True	

	mature	language	watch_time_hours	stream_time_hours	\
0	False	English	1.032694e+08	3587.50	
1	False	English	1.015280e+08	3530.75	
2	True	Portuguese	9.407652e+07	8588.00	

3	False	English	6.617197e+07	8629.00
4	False	English	6.118333e+07	2061.00

	followers_per_hour	views_per_follower	engagement_rate
0	483.570732	53.629351	0.008538
1	388.071656	65.470013	0.004823
2	119.210410	100.228279	0.006209
3	81.583729	151.348098	0.001955
4	1003.602135	38.192647	0.003312

```
[21]: # use all these features to predict watch_time_minutes
features = ['stream_time_minutes', 'average_viewers', 'peak_viewers',
            'followers', 'followers_gained', 'views_gained',
            'mature', 'views_per_follower', 'engagement_rate']

# language needs to be one-hot encoded as it is a categorical var
df_ml = pd.get_dummies(df, columns=['language'], drop_first=True)
df_ml.columns = [col.lower() if col.startswith('language_') else col for col in
                 df_ml.columns]

# filter for features and language columns
lang_columns = [col.lower() for col in df_ml.columns if col.
                 startswith('language_')]
X = df_ml[features + lang_columns].copy()
y = df_ml['watch_time_minutes']

# handle any missing vals
X = X.fillna(X.median())

# now we can perform a train/test split
# 80% for training and 20% for testing
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)

# scale the features so they each contribute equally
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

linear_model = LinearRegression()
linear_model.fit(X_train_scaled, y_train)

# evaluate linear regression model
y_pred_linear = linear_model.predict(X_test_scaled)
r2_linear = r2_score(y_test, y_pred_linear)
rmse_linear = np.sqrt(mean_squared_error(y_test, y_pred_linear))
mae_linear = mean_absolute_error(y_test, y_pred_linear)
```

```

print(f"Linear Regression Results:")
print(f"R2 Score: {r2_linear:.4f}")
print(f"RMSE: {rmse_linear:.2f}")
print(f"MAE: {mae_linear:.2f}")

```

Linear Regression Results:

R² Score: 0.6536

RMSE: 332317064.43

MAE: 163213427.06

```

[35]: def plot_actual_vs_predicted(model, X_test_scaled, y_test):
    y_pred = model.predict(X_test_scaled)

    plt.figure(figsize=(12, 8))
    plt.scatter(y_test, y_pred)
    plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], 'r--')
    plt.xlabel('Actual Watch Time (minutes)')
    plt.ylabel('Predicted Watch Time (minutes)')
    plt.title('Linear Regression: Actual vs Predicted Watch Time', fontsize=14)

    # Add performance metrics to plot
    r2 = r2_score(y_test, y_pred)
    rmse = np.sqrt(mean_squared_error(y_test, y_pred))
    plt.annotate(f'R2 = {r2:.3f}\nRMSE = {rmse:.0f}',
                xy=(0.05, 0.95), xycoords='axes fraction',
                bbox=dict(boxstyle="round,pad=0.3", fc="white", ec="gray",
                ↪alpha=0.8))

    plt.show()

```

```

[36]: # Option 2: Coefficient Plot (Feature Importance)
def plot_linear_coefficients(model, feature_names):
    # Get coefficients
    coefficients = pd.DataFrame({
        'Feature': feature_names,
        'Coefficient': model.coef_
    })

    # Sort by absolute value
    coefficients['Abs_Coefficient'] = np.abs(coefficients['Coefficient'])
    coefficients = coefficients.sort_values('Abs_Coefficient', ascending=False)

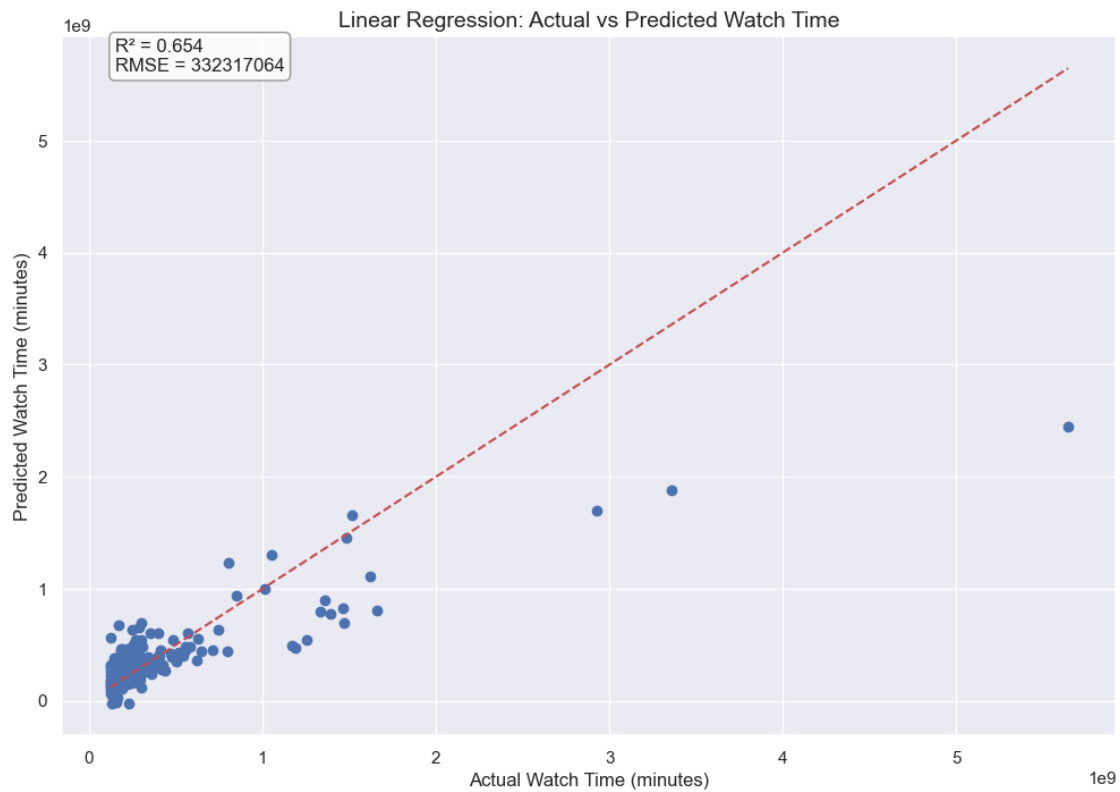
    # Plot top 15 coefficients
    plt.figure(figsize=(12, 8))

    sb.barplot(x='Coefficient', y='Feature', data=coefficients.head(15))

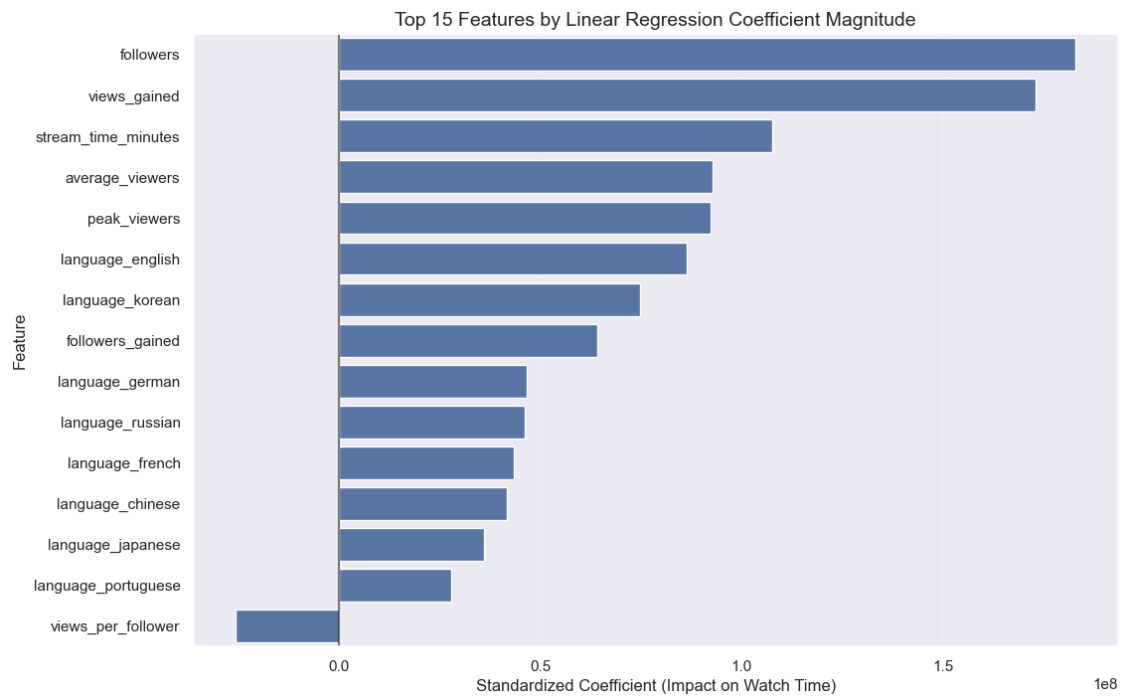
```

```
plt.title('Top 15 Features by Linear Regression Coefficient Magnitude',
↪fontsize=14)
plt.xlabel('Standardized Coefficient (Impact on Watch Time)', fontsize=12)
plt.axvline(x=0, color='black', linestyle='-', alpha=0.5)
plt.grid(axis='x', alpha=0.3)
plt.show()
```

```
[37]: plot_actual_vs_predicted(linear_model, X_test_scaled, y_test)
```



```
[38]: plot_linear_coefficients(linear_model, features + lang_columns)
```



[]: